


Long-context crossover (EC6)



Long-context crossover (EC6): acc-per-token (mean ACC / mean tokens, the preregistered metric) for treatment vs baseline across stream length L. no cost-adjusted crossover in the swept range. n=8 seeds; bands/error bars = 95% CI (normal approx, mean +/- 1.96*SE over seeds). MOCK-LLM CAVEAT: MECHANISM DEMONSTRATION ONLY. Every number in this analysis was produced in the synthetic fact-world stream under the deterministic mock LLM (no ANTHROPIC_API_KEY on this box), so it is reproducible bit-for-bit. The H1/H0 verdict is therefore a statement about the bench mechanism in this synthetic + mock-LLM regime, NOT a scientific claim about a real Claude model. Neuroscience (sleep, replay, TMR) is motivation and analogy only, never a proof of biological fidelity or real-model efficacy (DX5).