


Prune quality per arm (distractor_heavy regime)



Mechanism-level prune quality (precision / recall / F1) per arm in the primary regime, decoupled from downstream accuracy. n=8 seeds; bands/error bars = 95% CI (normal approx, mean $\pm 1.96 \cdot \text{SE}$ over seeds). **MOCK-LLM CAVEAT:** MECHANISM DEMONSTRATION ONLY. Every number in this analysis was produced in the synthetic fact-world stream under the deterministic mock LLM (no ANTHROPIC_API_KEY on this box), so it is reproducible bit-for-bit. The H1/H0 verdict is therefore a statement about the bench mechanism in this synthetic + mock-LLM regime, NOT a scientific claim about a real Claude model. Neuroscience (sleep, replay, TMR) is motivation and analogy only, never a proof of biological fidelity or real-model efficacy (DX5).