


Retention curves (distractor_heavy regime)



Per-task retention (final-row accuracy $R[T-1, j]$ after the full stream) per arm in the 'distractor_heavy' regime. $n=8$ seeds; bands/error bars = 95% CI (normal approx, mean $\pm 1.96 \cdot SE$ over seeds). **MOCK-LLM CAVEAT: MECHANISM DEMONSTRATION ONLY.** Every number in this analysis was produced in the synthetic fact-world stream under the deterministic mock LLM (no ANTHROPIC_API_KEY on this box), so it is reproducible bit-for-bit. The H1/H0 verdict is therefore a statement about the bench mechanism in this synthetic + mock-LLM regime, NOT a scientific claim about a real Claude model. Neuroscience (sleep, replay, TMR) is motivation and analogy only, never a proof of biological fidelity or real-model efficacy (DX5).