


TMR-style targeting effect (FR5.3)



TMR-style targeting: signal-retention lift for replay arms over no-replay arms (a replay-targeting analogue, NOT a literal cued-TMR protocol), vs the Hu et al. 2020 benchmark Hedges' $g=0.29$. pooled obs: replay=24, no_replay=16; $g=1.70$. $n=8$ seeds; bands/error bars = 95% CI (normal approx, mean $\pm 1.96 \cdot SE$ over seeds); Hedges' g error bar = bootstrap 95% CI. MOCK-LLM CAVEAT: MECHANISM DEMONSTRATION ONLY. Every number in this analysis was produced in the synthetic fact-world stream under the deterministic mock LLM (no ANTHROPIC_API_KEY on this box), so it is reproducible bit-for-bit. The H1/H0 verdict is therefore a statement about the bench mechanism in this synthetic + mock-LLM regime, NOT a scientific claim about a real Claude model. Neuroscience (sleep, replay, TMR) is motivation and analogy only, never a proof of biological fidelity or real-model efficacy (DX5).